

Mermade — Zama Merkle Tree Client/Server Challenge

Challenge

Imagine a client has a large set of potentially small files $\{F_0, F_1, \dots, F_n\}$ and wants to upload them to a server and then delete its local copies. The client wants, however, to later download an arbitrary file from the server and be convinced that the file is correct and is not corrupted in any way (in transport, tampered with by the server, etc.).

You should implement the client, the server and a Merkle tree to support the above (we expect you to implement the Merkle tree rather than use a library, but you are free to use a library for the underlying hash functions).

The client must compute a single Merkle tree root hash and keep it on its disk after uploading the files to the server and deleting its local copies. The client can request the i -th file F_i and a Merkle proof P_i for it from the server. The client uses the proof and compares the resulting root hash with the one it persisted before deleting the files - if they match, file is correct.

You can use any programming language you want (we use Rust internally). We would like to see a solution with networking that can be deployed across multiple machines, and as close to production-ready as you have time for. Please describe the short-comings your solution has in a report, and how you would improve on them given more time.

We expect you to send us within 7 days:

a demo of your app that we can try (ideally using eg Docker Compose) the code of the app a report (max 2-3 pages) explaining your approach, your other ideas, what went well or not, etc..

Assumptions

1. I assume the purpose of the challenge is to demonstrate the ability to code, design and implement a solution that is production ready, but not necessarily to implement a production ready solution. I will therefore make some assumptions and take some shortcuts to save our time.
2. There are many questions that I would ask if this was a real project, but I won't. I'll make more assumptions instead.
3. I assume the "large set of potentially small files" means that each file is small enough to fit in memory to calculate its hash without streaming, and "large set" means up to several millions of files. My implementation stores the whole Merkle tree in memory, requiring ~64 bytes per file. So, for 1

million files it will require ~64MB of memory. I assume this is acceptable on a modern client or server.

4. I assume the files to upload are named by their index, like “0”, “1”, “2”, etc.
5. I assume the client and the server are on the same network, so I don’t need to implement any authentication or encryption.
6. I assume the server has enough disk space to store all files.
7. I assume the server has only one client, although it’s easy to extend the solution to support multiple clients, and even shard to multiple servers using a load balancer.

Solution

I’ve implemented a solution in Rust, using Actix Web framework for the server and Reqwest for the client.

In prod-level solution I would split the project into 3 separate sub-projects: shared library, client and server. Here, I’ll keep everything in one application that can be run as a client or a server. My approach simplifies code review, which is the purpose of the challenge.

The client reads the current working directory.

It calculates the Merkle tree for all files in the directory and stores the Merkle Root in `$HOME` directory.

The server IP address is passed as a command line argument.

Then it uploads all files to the server, one by one, using REST API POST “/upload”. I could use multipart upload, but I decided to keep it simple.

A client can request a file by its index using REST API GET “/files/{index}”.

A client can request a proof for a file by its index using REST API GET “/proofs/{index}”.

The client verifies the file using the Merkle proof and the Merkle Root it has stored in `$HOME` directory.

The server stores all files in a directory named “files” in the current working directory.

On client’s first GET request, the server computes Merkle proofs for each file and stores them in *proof* files in “proofs” directory.

This is a very simple and efficient solution. The Merkle tree and proofs are computed only once, and proofs are essentially cached. Serving static files is very efficient.

Then, on client GET request, the server would simply `sendfile` the file and the proof file to the client. It would be very fast and efficient.

Merkle Tree

I compute and store the Merkle tree in memory in the following form on both client and server. It's a vector of levels of the tree, where each level is a vector of hashes of the nodes on that level. So, for 3 files with hashes "aa", "bb", "cc" the tree will look like this:

```
[
  ["aa", "bb", "cc", "cc"],
  ["dd", "ee"],
  ["ff"]
]
```

where "ff" is the Merkle Root.

This is not the most efficient way to store the tree, but it's simple, easy to implement, and it works. From this implementation it's trivial to derive both Merkle root and proofs.

If needed I can implement a "rolling" Merkle root computation, requiring $\sim 2 \cdot \log_2(N)$ memory, where N is the number of files.

Other

In prod-level solution I would add logging, metrics, more tests, and better error handling, configuration, multiple clients, file sharing, merkle proofs recalculation on file changes, backpressure, rate limiting, etc.

I've implemented this in Rust mostly for fun, and because you use Rust internally. I'm not a Rust expert, so I'm sure there are many things I could do better.